AbroadShield AI

Consolidated Project Brief

Internal reference — every locked decision to date, in one place

July 24, 2026

1. The Core Idea

AbroadShield AI is a single, continuous, agentic AI that walks a student through the entire journey abroad — not four separate tools, one relationship with one memory, spanning four phases: Pre-Departure, Arrival, Studying & Part-Time, and Job Success.

It is agentic, not conversational: it does not wait to be asked. With approval, it drafts emails and forms, searches and shortlists housing and jobs, checks documents for gaps before appointments, and tracks every deadline — acting instead of only advising.

It is built for the full size of the real market — including tier-2 and tier-3 Indian towns, not only metro, English-fluent students — and it stays current on the country-specific rules (registration windows, insurance mandates, work-hour limits) that rarely appear on a generic checklist.

2. Why This Beats Using Claude, Grok, or Gemini Directly

This question was pressure-tested deliberately, across multiple models, specifically to avoid an echo chamber. The answer that survived scrutiny rests on four pillars:

- Persistent, structured memory of one specific journey — not memory in general, but memory organized around a real sequence (visa → flight → SIM → housing → bank → work-hour limits) for a specific country.
- Proactive outbound nudges — the product reaches out first, before a deadline, rather than waiting to be asked.
- Domain-specific journey mapping — it already knows the checklist, sequence, and country rules; the student never has to explain their situation from scratch.
- Agentic task execution — the strongest pillar. It drafts, searches, shortlists, and fills in, with approval, rather than only explaining what to do. This is the one general AI assistants will not do for a user by default, because it requires purpose-built automation most people will never configure themselves.

Important correction carried forward: general AI assistants (Claude, ChatGPT, Gemini) do have real persistent memory today. The accurate differentiation is not "they forget everything" — it is that they are general-purpose, reactive by default, and not structured around one specific country's rules or one student's specific deadlines.

3. Moat & Business Risk Analysis

Research into AI "wrapper" products confirms the real risk: the large majority of thin AI-wrapper products fail, and foundation labs can and do absorb undefended features into their default products overnight ("getting Sherlocked"). The test applied: if a foundation model shipped this exact pitch as a default feature, would users cancel? The honest answer for AbroadShield is not immediately — because of the moats below.

Moat	**Current Status**
Distribution	Real, and the strongest one today — reaching underserved Indian students in tier-2/3 towns, in regional languages, is a segment large AI labs have no reason to build for specifically.
Workflow / memory lock-in	Plausible, not yet proven — becomes real once a student's full multi-year history lives inside the product, making switching costly.
Data moat	Does not exist yet — must be built deliberately by tracking real outcomes (which documents led to visa rejections, which country rules changed when).
Compliance / regulatory	Thin but real — staying accurately current on country-specific visa and settlement rules is closer to a compliance moat than a chatbot feature.
Brand / trust	Earned over time — not present yet, builds with consistent delivery.

Verdict: build the boring infrastructure first — the state machine and outbound nudge engine — not the "Godfather" personality. The personality is the easiest part to copy; the sequencing and tracking engine is the hard, defensible part.

4. Technical Architecture (Core, Locked)

- Backend: FastAPI, handling inbound WhatsApp webhooks (via AiSensy) and driving agent logic.

- State & data: Supabase (Postgres) — student profiles, milestone/state tracking, document vault, outbound nudge log.
- Agent orchestration: LangGraph-based state machine for milestone sequencing.
- LLM: Claude or Grok API for conversational and drafting tasks.
- Added table (identified gap): a milestone-template table mapping destination country → phase → required steps — this is what auto-generates a student's own checklist, and is arguably the most valuable proprietary data the product owns.
- Two code-level fixes applied to the original backend skeleton: Pydantic model must inherit from BaseModel (not a raw dict param), and route decorators must be called on the app instance (@app.get, not a bare @get).

Distribution partnerships (e.g. consultants, coaching centers) were discussed as a possible future channel to solve cold-start discovery, but this is explicitly deferred — not part of the current build focus.

5. Business Model (Final)

Consultant/agency licensing was proposed twice and explicitly rejected both times — it is not part of the model. The current locked model is direct-to-student, with future partnerships kept separate from near-term revenue.

Revenue Stream	Notes
Direct student subscription	Free tier (checklist + nudges) and a paid tier unlocking agentic actions, from day one — not gated behind a trust-building period.
Affiliate commissions	Insurance, forex, SIM, housing — real transactions the student makes anyway, no added cost to them.
Premium Job-Success tier	Highest willingness-to-pay moment in the journey (Phase 4).
Future partnerships	Universities (retention pipeline), financial services (embedded banking/forex/insurance), employers & recruiters (visa-ready candidate pool), government/embassy-adjacent programs, diaspora & alumni networks. Explicitly not pursued yet — noted for later, not part of near-term revenue.

6. New Addition: Networking & Alumni/HR Outreach (Phase 4)

A newly agreed enhancement to the Job Success phase: today, succeeding abroad increasingly depends on proactive networking — reaching out to alumni, seniors, local event contacts, and HR/hiring staff — not just submitting applications. The agent's role:

- Maintains a structured, always-on networking tracker — who has been contacted, when, what was said, and when to follow up — the same kind of tracking discipline as the deadline engine in Phase 1.
- Finds and surfaces relevant public events — alumni meetups, local industry meetups — via public event listings and university alumni pages.
- Drafts personalized outreach messages to alumni, seniors, or HR contacts, ready for the student's review.

Known constraint, consistent with the earlier LinkedIn decision: there is no public API for automating actions on LinkedIn, and doing so risks the student's own account being restricted. The agent researches, tracks, and drafts — the student sends the final message themselves. This preserves the full value of the capability without the platform risk.

7. Validation Survey

A 13-question Google Form was built and distributed to test the core assumptions directly with real students and families, covering all three groups (already abroad, currently planning, and just exploring) with a single consistent set of questions. Structure: mostly tap-to-select with an “Other” write-in option on every relevant question, one open free-text question reserved for “anything you wish someone had told you,” and Google Forms' built-in email collection and response-limiting enabled.

Distribution approach: direct personal outreach first (highest-quality responses), then wider channels such as relevant Reddit communities — approached carefully, respecting each subreddit's self-promotion and survey rules, framed as genuine curiosity rather than promotion, with realistic expectations (dozens of thoughtful responses, not guaranteed virality).

8. Immediate Next Steps

- Collect and review survey responses — let real answers, not assumptions, determine which phase and which country rules to build first.
- Finalize the Supabase schema, including the milestone-template table.
- Build the Core: state tracking, document vault, and the outbound nudge engine — before any personality or chat polish.
- Build Phase 1 (Pre-Departure) only, test it with real respondents from the survey, then expand phase by phase.